

```
[acadgild@localhost ~]$ cd Desktop
[acadgild@localhost Desktop]$ ls
gnome-terminal.desktop  hadoop-mapreduce-client-core-2.9.0.jar  new file
hadoop-common-2.9.0.jar  hadoop-mapreduce-examples-2.9.0.jar    wordcount.zip
[acadgild@localhost Desktop]$ hadoop fs -cat file1.txt
18/12/04 13:54:09 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your
platform... using builtin-java classes where applicable
cat: Call From localhost/127.0.0.1 to localhost:9000 failed on connection exception:
java.net.ConnectException: Connection refused; For more details see:
http://wiki.apache.org/hadoop/ConnectionRefused
[acadgild@localhost Desktop]$ start-all.sh
This script is Deprecated. Instead use start-dfs.sh and start-yarn.sh
18/12/04 13:54:22 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your
platform... using builtin-java classes where applicable
Starting namenodes on [localhost]
localhost: starting namenode, logging to /usr/local/hadoop-2.6.0/logs/hadoop-acadgild-namenode-
localhost.localdomain.out
localhost: starting datanode, logging to /usr/local/hadoop-2.6.0/logs/hadoop-acadgild-datanode-
localhost.localdomain.out
Starting secondary namenodes [0.0.0.0]
0.0.0.0: starting secondarynamenode, logging to /usr/local/hadoop-2.6.0/logs/hadoop-acadgild-
secondarynamenode-localhost.localdomain.out
18/12/04 13:54:58 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your
platform... using builtin-java classes where applicable
starting yarn daemons
starting resourcemanager, logging to /usr/local/hadoop-2.6.0/logs/yarn-acadgild-resourcemanager-
localhost.localdomain.out
localhost: starting nodemanager, logging to /usr/local/hadoop-2.6.0/logs/yarn-acadgild-
nodemanager-localhost.localdomain.out
[acadgild@localhost Desktop]$ hadoop fs -cat file1.txt
18/12/04 13:55:21 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your
platform... using builtin-java classes where applicable
A B C
C A B
B B A
B C A
A C C
```

Word Median

```
[acadgild@localhost Desktop]$ hadoop jar hadoop-mapreduce-examples-2.9.0.jar wordmedian
file1.txt wmedian
18/12/04 13:56:36 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your
platform... using builtin-java classes where applicable
18/12/04 13:56:38 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
18/12/04 13:56:40 WARN mapreduce.JobSubmitter: Hadoop command-line option parsing not
performed. Implement the Tool interface and execute your application with ToolRunner to remedy
this.
18/12/04 13:56:42 INFO input.FileInputFormat: Total input paths to process : 1
18/12/04 13:56:42 INFO mapreduce.JobSubmitter: number of splits:1
```

18/12/04 13:56:42 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1543911908873_0001
18/12/04 13:56:43 INFO impl.YarnClientImpl: Submitted application
application_1543911908873_0001
18/12/04 13:56:43 INFO mapreduce.Job: The url to track the job:
http://localhost:8088/proxy/application_1543911908873_0001/
18/12/04 13:56:43 INFO mapreduce.Job: Running job: job_1543911908873_0001
18/12/04 13:57:08 INFO mapreduce.Job: Job job_1543911908873_0001 running in uber mode : false
18/12/04 13:57:08 INFO mapreduce.Job: map 0% reduce 0%
18/12/04 13:57:20 INFO mapreduce.Job: map 100% reduce 0%
18/12/04 13:57:34 INFO mapreduce.Job: map 100% reduce 100%
18/12/04 13:57:35 INFO mapreduce.Job: Job job_1543911908873_0001 completed successfully
18/12/04 13:57:35 INFO mapreduce.Job: Counters: 49

File System Counters

FILE: Number of bytes read=16
FILE: Number of bytes written=211897
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=141
HDFS: Number of bytes written=5
HDFS: Number of read operations=6
HDFS: Number of large read operations=0
HDFS: Number of write operations=2

Job Counters

Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=9858
Total time spent by all reduces in occupied slots (ms)=10203
Total time spent by all map tasks (ms)=9858
Total time spent by all reduce tasks (ms)=10203
Total vcore-seconds taken by all map tasks=9858
Total vcore-seconds taken by all reduce tasks=10203
Total megabyte-seconds taken by all map tasks=10094592
Total megabyte-seconds taken by all reduce tasks=10447872

Map-Reduce Framework

Map input records=5
Map output records=15
Map output bytes=120
Map output materialized bytes=16
Input split bytes=110
Combine input records=15
Combine output records=1
Reduce input groups=1
Reduce shuffle bytes=16
Reduce input records=1
Reduce output records=1
Spilled Records=2
Shuffled Maps =1
Failed Shuffles=0

```

Merged Map outputs=1
GC time elapsed (ms)=367
CPU time spent (ms)=2360
Physical memory (bytes) snapshot=294232064
Virtual memory (bytes) snapshot=4113776640
Total committed heap usage (bytes)=165810176

Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0

File Input Format Counters
    Bytes Read=31

File Output Format Counters
    Bytes Written=5

The median is: 1
[acadgild@localhost Desktop]$ hadoop fs -ls wmedian
18/12/04 13:58:30 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your
platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r--  1 acadgild supergroup      0 2018-12-04 13:57 wmedian/_SUCCESS
-rw-r--r--  1 acadgild supergroup      5 2018-12-04 13:57 wmedian/part-r-000000
[acadgild@localhost Desktop]$ hadoop fs -cat wmedian/part-r-000000
18/12/04 13:59:13 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your
platform... using builtin-java classes where applicable
1          15

```

Word Mean

```

[acadgild@localhost Desktop]$ hadoop jar hadoop-mapreduce-examples-2.9.0.jar wordmean
file1.txt wmean
18/12/04 14:00:00 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your
platform... using builtin-java classes where applicable
18/12/04 14:00:02 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
18/12/04 14:00:05 INFO input.FileInputFormat: Total input paths to process : 1
18/12/04 14:00:05 INFO mapreduce.JobSubmitter: number of splits:1
18/12/04 14:00:06 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1543911908873_0002
18/12/04 14:00:06 INFO impl.YarnClientImpl: Submitted application
application_1543911908873_0002
18/12/04 14:00:07 INFO mapreduce.Job: The url to track the job:
http://localhost:8088/proxy/application_1543911908873_0002/
18/12/04 14:00:07 INFO mapreduce.Job: Running job: job_1543911908873_0002
18/12/04 14:00:22 INFO mapreduce.Job: Job job_1543911908873_0002 running in uber mode : false
18/12/04 14:00:22 INFO mapreduce.Job: map 0% reduce 0%
18/12/04 14:00:33 INFO mapreduce.Job: map 100% reduce 0%
18/12/04 14:00:46 INFO mapreduce.Job: map 100% reduce 100%
18/12/04 14:00:46 INFO mapreduce.Job: Job job_1543911908873_0002 completed successfully
18/12/04 14:00:46 INFO mapreduce.Job: Counters: 49

```

File System Counters

FILE: Number of bytes read=39
FILE: Number of bytes written=212203
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=141
HDFS: Number of bytes written=19
HDFS: Number of read operations=6
HDFS: Number of large read operations=0
HDFS: Number of write operations=2

Job Counters

Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=8739
Total time spent by all reduces in occupied slots (ms)=9276
Total time spent by all map tasks (ms)=8739
Total time spent by all reduce tasks (ms)=9276
Total vcore-seconds taken by all map tasks=8739
Total vcore-seconds taken by all reduce tasks=9276
Total megabyte-seconds taken by all map tasks=8948736
Total megabyte-seconds taken by all reduce tasks=9498624

Map-Reduce Framework

Map input records=5
Map output records=30
Map output bytes=435
Map output materialized bytes=39
Input split bytes=110
Combine input records=30
Combine output records=2
Reduce input groups=2
Reduce shuffle bytes=39
Reduce input records=2
Reduce output records=2
Spilled Records=4
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=328
CPU time spent (ms)=2220
Physical memory (bytes) snapshot=293777408
Virtual memory (bytes) snapshot=4113768448
Total committed heap usage (bytes)=165810176

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=31

File Output Format Counters

Bytes Written=19

The mean is: 1.0

[acadgild@localhost Desktop]\$ hadoop fs -ls wmean

18/12/04 14:01:07 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 2 items

-rw-r--r-- 1 acadgild supergroup 0 2018-12-04 14:00 wmean/_SUCCESS

-rw-r--r-- 1 acadgild supergroup 19 2018-12-04 14:00 wmean/part-r-00000

[acadgild@localhost Desktop]\$ hadoop fs -cat wmean/part-r-00000

18/12/04 14:01:29 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

count 15

length 15

Word Standard deviation

[acadgild@localhost Desktop]\$ hadoop jar hadoop-mapreduce-examples-2.9.0.jar

wordstandarddeviation file1.txt wstd

18/12/04 14:02:40 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

18/12/04 14:02:41 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032

18/12/04 14:02:44 INFO input.FileInputFormat: Total input paths to process : 1

18/12/04 14:02:44 INFO mapreduce.JobSubmitter: number of splits:1

18/12/04 14:02:45 INFO mapreduce.JobSubmitter: Submitting tokens for job:

job_1543911908873_0003

18/12/04 14:02:45 INFO impl.YarnClientImpl: Submitted application

application_1543911908873_0003

18/12/04 14:02:45 INFO mapreduce.Job: The url to track the job:

http://localhost:8088/proxy/application_1543911908873_0003/

18/12/04 14:02:45 INFO mapreduce.Job: Running job: job_1543911908873_0003

18/12/04 14:03:01 INFO mapreduce.Job: Job job_1543911908873_0003 running in uber mode : false

18/12/04 14:03:01 INFO mapreduce.Job: map 0% reduce 0%

18/12/04 14:03:12 INFO mapreduce.Job: map 100% reduce 0%

18/12/04 14:03:24 INFO mapreduce.Job: map 100% reduce 100%

18/12/04 14:03:26 INFO mapreduce.Job: Job job_1543911908873_0003 completed successfully

18/12/04 14:03:26 INFO mapreduce.Job: Counters: 49

File System Counters

FILE: Number of bytes read=56

FILE: Number of bytes written=212395

FILE: Number of read operations=0

FILE: Number of large read operations=0

FILE: Number of write operations=0

HDFS: Number of bytes read=141

HDFS: Number of bytes written=29

HDFS: Number of read operations=6

HDFS: Number of large read operations=0

HDFS: Number of write operations=2

Job Counters

Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=8492
Total time spent by all reduces in occupied slots (ms)=9304
Total time spent by all map tasks (ms)=8492
Total time spent by all reduce tasks (ms)=9304
Total vcore-seconds taken by all map tasks=8492
Total vcore-seconds taken by all reduce tasks=9304
Total megabyte-seconds taken by all map tasks=8695808
Total megabyte-seconds taken by all reduce tasks=9527296

Map-Reduce Framework

Map input records=5
Map output records=45
Map output bytes=660
Map output materialized bytes=56
Input split bytes=110
Combine input records=45
Combine output records=3
Reduce input groups=3
Reduce shuffle bytes=56
Reduce input records=3
Reduce output records=3
Spilled Records=6
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=408
CPU time spent (ms)=2220
Physical memory (bytes) snapshot=294883328
Virtual memory (bytes) snapshot=4115869696
Total committed heap usage (bytes)=165810176

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=31

File Output Format Counters

Bytes Written=29

The standard deviation is: 0.0

[acadgild@localhost Desktop]\$ hadoop fs -ls wstd

18/12/04 14:03:44 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 2 items

```
-rw-r--r-- 1 acadgild supergroup      0 2018-12-04 14:03 wstd/_SUCCESS
-rw-r--r-- 1 acadgild supergroup    29 2018-12-04 14:03 wstd/part-r-00000
[acadgild@localhost Desktop]$ hadoop fs -cat wstd/part-r-00000
```

```
18/12/04 14:04:07 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your
platform... using builtin-java classes where applicable
count  15
length 15
square 15
[acadgild@localhost Desktop]$
```